

# devi v0.4.0 - Comprehensive Performance Characterization

This release contains three independent performance studies on parallel chess search:

## 1. Heterogeneous Core Scheduling

- Problem: Can we exploit Apple Silicon's P/E asymmetry for chess search?
- Method: QoS-based thread biasing across 4 policies
- Result: E-cores **12.8x slower**, mixed policy achieves only 65% of expected
- Details: [core\\_pinning\\_analysis.md](#)

## 2. Scaling Laws (Amdahl vs Gustafson)

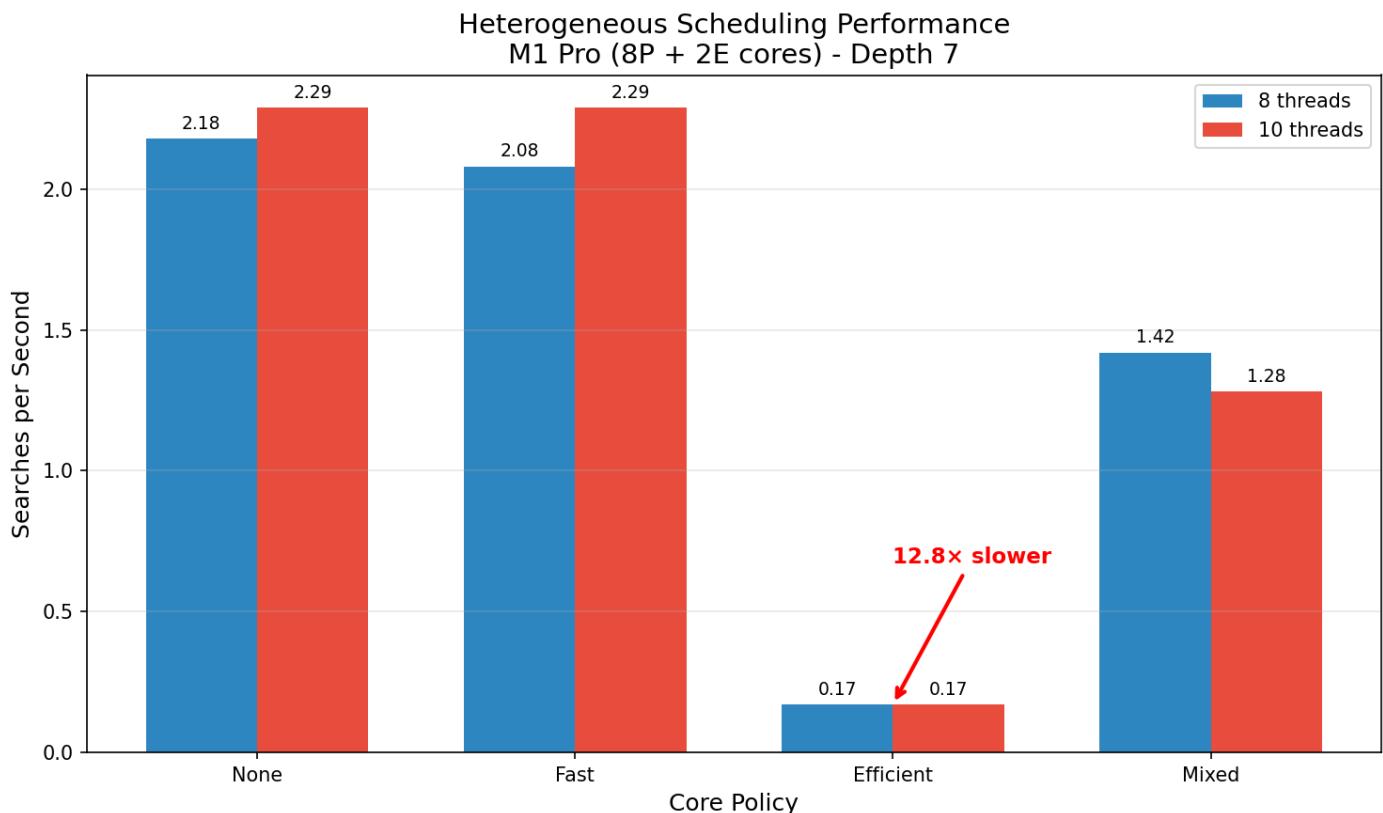
- Problem: Does chess search follow Amdahl's Law (fixed problem) or Gustafson's Law (scaled problem)?
- Method: Compare scaling at depth 4 vs depth 7
- Result: 48% serial fraction reduction (0.27 > 0.14) as depth increases
- Details: [scaling\\_analysis.md](#)

## 3. Fault Tolerance Baseline

- Problem: What's the cost of thread-level fault recovery without checkpointing?
- Method: Inject panic at move 5 after 2-ply work, measure full retry overhead
- Result: **100% overhead** - Rayon's all-or-nothing model discards all parallel work
- Details: [fault\\_tolerance\\_analysis.md](#)

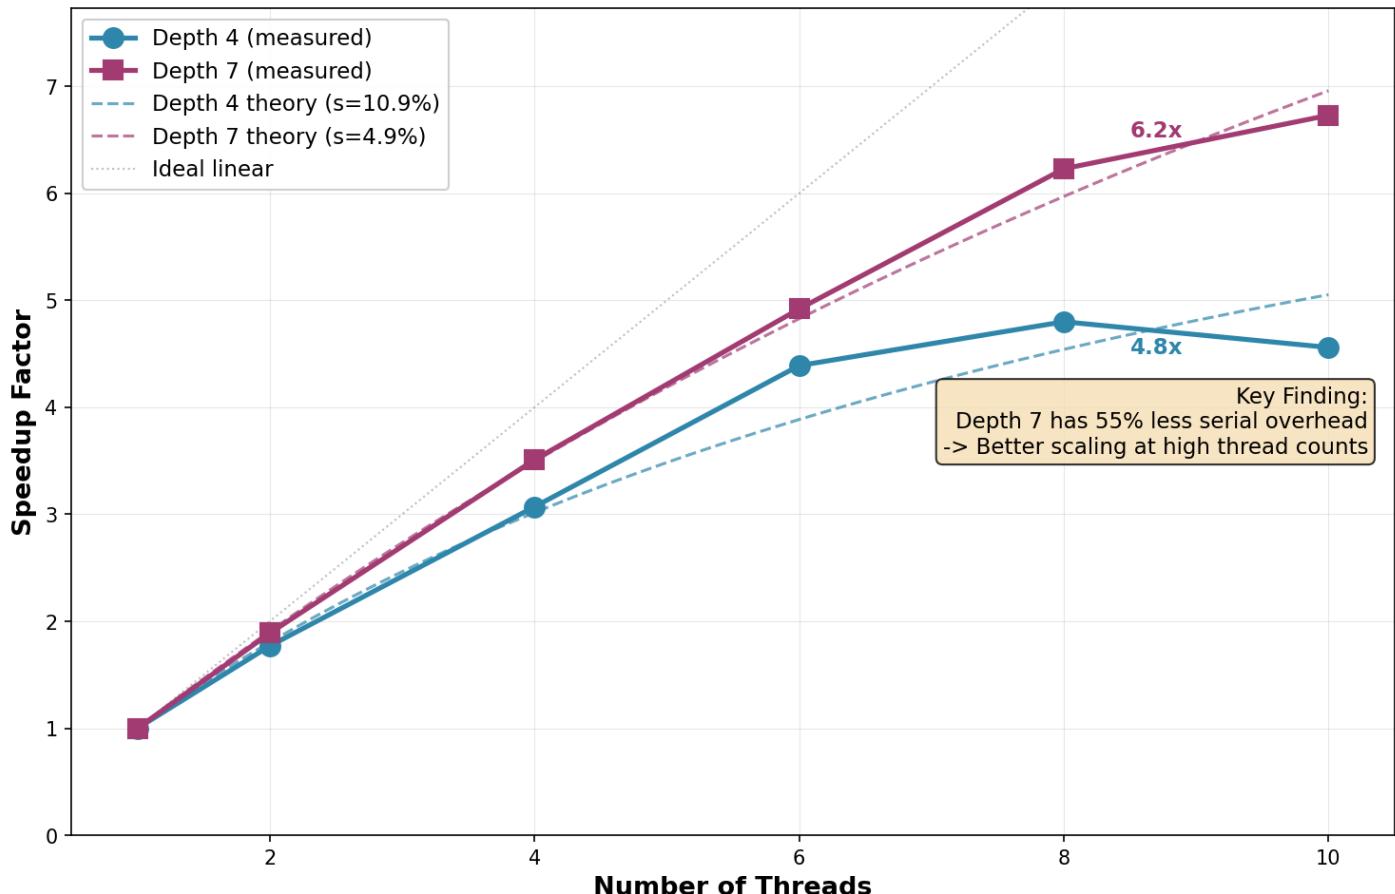
## Summary

### Heterogeneous Performance



## Scaling Comparison

### Parallel Speedup: Small Problem (Depth 4) vs Large Problem (Depth 7)



## Key Findings

1. Fault recovery needs checkpointing - 100% overhead unacceptable for production
2. Larger problems scale better - Gustafson's Law dominates at higher depths
3. E-cores unsuitable for branch-heavy work - 12.8x gap motivates separate work queues
  - Apple M1 Pro (8P + 2E cores)

## Reproduce All Studies

```
# Clone and build
git clone https://github.com/Sid4mn/devi-chess-engine.git
cd chess-engine-rust && cargo build --release

# 1. Fault tolerance
./target/release/devi --fault-analysis --depth 7 --threads 8

# 2. Scaling laws
./scripts/analysis/multi_depth_scaling.sh

# 3. Heterogeneous scheduling
./scripts/heterogeneous.sh
```

## Data Files

- Fault tolerance: [benchmarks/fault\\_overhead.csv](#)
- Heterogeneous: [benchmarks/hetero\\_\\*.csv](#)
- Scaling: [benchmarks/scaling\\_analysis/](#)

## Future Work: Harness-Inspired Lightweight Scheduler

---

1. Our two-phase root execution design adapts concepts from Belviranli et al.'s Harness framework [1], which demonstrated that **pattern-based work classification** can effectively utilize heterogeneous cores. While Harness targets MapReduce workloads with runtime pattern detection, we apply similar principles to chess search:

- **Harness approach:** Detect map/reduce patterns -> route to appropriate cores
- **Our adaptation:** PV probe estimates subtree complexity -> route heavy/light moves to P/E pools

The key insight from Harness is that **workload characterization enables better heterogeneous scheduling** which directly motivates our PV probe design.

[1] Belviranli, M. E., et al. "Harness: A Pattern-based Framework for Optimizing Heterogeneous MapReduce Applications." ICS 2025.

2. **Checkpoint Recovery** - Per-thread root-level checkpointing with `catch_unwind` to preserve completed move evaluations on panic. Target:  $\leq 30\%$  overhead (vs current 100% full retry).
-